

THE CORE 50 EQUATIONS OF ARTIFICIAL INTELLIGENCE

The 50 mathematical ideas that power modern AI

I. LEARNING (10)	II. REPRESENTATION (10)	III. COMPUTATION (10)	IV. OPTIMISATION (10)	V. THEORY (10)
1 Unknown Function $f(x)$	11 Representation Function $\phi: X \rightarrow Z$	21 Linear Neuron $y = w^T x + b$	31 Loss Function $L(y, \hat{y})$	41 VC-style Intuition (Bound) $R \leq \hat{R} + \text{Complexity}$
2 Parameterised Function $f_\theta(x)$	12 Latent Representation $z = \phi(x)$	22 Activation $a = \sigma(w^T x + b)$	32 Gradient ∇L	42 PAC Learning $P(R(f) \leq \epsilon) \geq 1 - \delta$
3 Input-Output Mapping $f: X \rightarrow Y$	13 Distance $d(x, y) = \ x - y\ $	23 Layer (Affine + Act.) $h = \sigma(Wx + b)$	33 Gradient Descent $\theta = \theta - \eta \nabla L$	43 Scaling Law $L(N) = AN^{-\alpha}$
4 Empirical Risk $\hat{R}(f) = \frac{1}{n} \sum_{i=1}^n L(y_i, f(x_i))$	14 Cosine Similarity $\frac{x \cdot y}{\ x\ \ y\ }$	24 Composition $f = g \circ h$	34 Chain Rule $\frac{dz}{dx} = \frac{dz}{dy} \frac{dy}{dx}$	44 Power Law $y = ax^b$
5 Expected Risk $R(f) = \mathbb{E}_{(x,y) \sim P} [L(Y, f(X))]$	15 Dot Product $x^T y$	25 Matrix Multiplication $Y = WX$	35 Mean Squared Error $\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$	45 Gaussian Distribution $\mathcal{N}(\mu, \sigma^2)$
6 Learning Objective $\theta^* = \underset{\theta}{\operatorname{argmin}} L(\theta)$	16 Principal Components $X = USV^T$	26 Convolution $(f * g)(x)$	36 Cross Entropy $-\sum_i y_i \log \hat{y}_i$	46 Expectation $\mathbb{E}[X]$
7 Generalisation Gap $R(f) - \hat{R}(f)$	17 Information Bottleneck $\max_{\phi} I(Z; Y) - \beta I(Z; X)$	27 Attention Scores QK^T	37 L2 Regularisation $L + \lambda \ w\ _2^2$	47 Variance $\operatorname{Var}(X) = \mathbb{E}[(X - \mu)^2]$
8 Maximum Likelihood $\theta^* = \underset{\theta}{\operatorname{argmax}} P(D \theta)$	18 Entropy $H(X) = -\sum_x p(x) \log p(x)$	28 Scaled Attention $\operatorname{softmax}\left(\frac{QK^T}{\sqrt{d}}\right) V$	38 L1 Regularisation $L + \lambda \ w\ _1$	48 Markov Property $P(S_{t+1} S_t)$
9 Bayes Rule $P(A B) = \frac{P(B A)P(A)}{P(B)}$	19 Mutual Information $I(X; Y) = H(X) - H(X Y)$	29 Softmax $P_i = \frac{e^{z_i}}{\sum_j e^{z_j}}$	39 Adam Update $\theta = \theta - \eta \frac{\hat{m}}{\sqrt{\hat{v}} + \epsilon}$	49 Bellman Equation $V(s) = R + \gamma V(s')$
10 Posterior $P(\theta D) \propto P(D \theta)P(\theta)$	20 KL Divergence $D_{KL}(P Q) = \sum_x P(x) \log \frac{P(x)}{Q(x)}$	30 Residual Connection $y = x + F(x)$	40 Bias-Variance Trade-off Error = Bias ² + Variance + Noise	50 Universal Approximation $f(x) \approx \sum_i a_i \sigma(w^T x + b_i)$

THE AI PIPELINE (BRINGING IT ALL TOGETHER)


 Data
 $(x, y) \sim P(x, y)$


 Representation
 $z = \phi(x)$


 Computation
 $\hat{y} = f_\theta(z)$


 Optimisation
 $\theta^* = \underset{\theta}{\operatorname{argmin}} L(\theta)$


 Theory
 Why it works!

HOW TO USE THESE EQUATIONS

- 📖 These 50 equations form the mathematical foundation of modern AI.
- 🔗 They reappear across models, algorithms and theories.
- 🧠 Master these, and you can understand (and build) almost anything in AI.
- 🔄 Revisit them often – depth comes with repetition and application.

★ Remember: This is not just a list – it is the language of AI. | Learn the language. See the connections. Build the future.

THE CORE 50 EQUATIONS OF ARTIFICIAL INTELLIGENCE

An Annotated Companion

For each equation: its plain-English meaning, the intuition beneath the notation, where it shows up in real systems, where it came from, what people get wrong about it, and why it still matters.

Five movements · Learning · Representation · Computation · Optimisation · Theory

By

[Ajit Jaokar](#)

For

The course [Mathematical foundations of Artificial Intelligence](#)

How to read this	6
Part I — Learning	7
01 Unknown Function	7
02 Parameterised Function	7
03 Input–Output Mapping	7
04 Empirical Risk	8
05 Expected Risk	8
06 Learning Objective	9
07 Generalisation Gap	9
08 Maximum Likelihood	9
09 Bayes Rule	10
10 Posterior	10
Part II — Representation	11
11 Representation Function	11
12 Latent Representation	11
13 Euclidean Distance	11
14 Cosine Similarity	12
15 Dot Product	12
16 Principal Components (SVD)	12
17 Information Bottleneck	13
18 Entropy	13
19 Mutual Information	14
20 KL Divergence	14
Part III — Computation	15
21 Linear Neuron	15
22 Activation	15
23 Layer	15
24 Composition	16
25 Matrix Multiplication	16
26 Convolution	16
27 Attention Scores	17
28 Scaled Dot-Product Attention	17
29 Softmax	17
30 Residual Connection	18
Part IV — Optimisation	19
31 Loss	19
32 Gradient	19
33 Gradient Descent	19
34 Chain Rule	20
35 Mean Squared Error	20
36 Cross Entropy	20
37 L2 Regularisation	21
38 L1 Regularisation	21
39 Adam Update	21

40 Bias–Variance	22
Part V — Theory	23
41 Generalisation Bound	23
42 PAC Learning	23
43 Scaling Law	23
44 Power Law	24
45 Gaussian	24
46 Expectation	25
47 Variance	25
48 Markov Property	25
49 Bellman Equation	26
50 Universal Approximation	26

Contents

Part I — Learning

- 01 Unknown Function $f(x)$
- 02 Parameterised Function $f\theta(x)$
- 03 Input–Output Mapping $f: X \rightarrow Y$
- 04 Empirical Risk $\hat{R}(f) = (1/n) \sum L(y_i, f(x_i))$
- 05 Expected Risk $R(f) = E[L(Y, f(X))]$
- 06 Learning Objective $\theta^* = \arg \min L(\theta)$
- 07 Generalisation Gap $R(f) - \hat{R}(f)$
- 08 Maximum Likelihood $\theta^* = \arg \max P(D|\theta)$
- 09 Bayes Rule $P(A|B) = P(B|A)P(A)/P(B)$
- 10 Posterior $P(\theta|D) \propto P(D|\theta)P(\theta)$

Part II — Representation

- 11 Representation Function $\varphi: X \rightarrow Z$
- 12 Latent Representation $z = \varphi(x)$
- 13 Euclidean Distance $d(x, y) = \|x - y\|$
- 14 Cosine Similarity $(x \cdot y) / (\|x\| \|y\|)$
- 15 Dot Product $x^T y$
- 16 Principal Components (SVD) $X = USV^T$
- 17 Information Bottleneck $\max I(Z; Y) - \beta I(Z; X)$
- 18 Entropy $H(X) = -\sum p(x) \log p(x)$
- 19 Mutual Information $I(X; Y) = H(X) - H(X|Y)$
- 20 KL Divergence $D_{KL}(P||Q) = \sum P(x) \log(P(x)/Q(x))$

Part III — Computation

- 21 Linear Neuron $y = w^T x + b$
- 22 Activation $a = \sigma(w^T x + b)$
- 23 Layer $h = \sigma(Wx + b)$
- 24 Composition $f = g \circ h$
- 25 Matrix Multiplication $Y = WX$
- 26 Convolution $(f * g)(x)$
- 27 Attention Scores QK^T

- 28 Scaled Dot-Product Attention $\text{softmax}((QK^T)/\sqrt{d})V$
- 29 Softmax $P_i = e^{z_i} / \sum e^{z_j}$
- 30 Residual Connection $y = x + F(x)$

Part IV — Optimisation

- 31 Loss $L(y, \hat{y})$
- 32 Gradient ∇L
- 33 Gradient Descent $\theta \leftarrow \theta - \eta \nabla L$
- 34 Chain Rule $dz/dx = (dz/dy)(dy/dx)$
- 35 Mean Squared Error $(1/n) \sum (y - \hat{y})^2$
- 36 Cross Entropy $-\sum y \log \hat{y}$
- 37 L2 Regularisation $L + \lambda \|w\|^2$
- 38 L1 Regularisation $L + \lambda \|w\|_1$
- 39 Adam Update $\theta \leftarrow \theta - \eta \hat{m}/(\sqrt{\hat{v}} + \epsilon)$
- 40 Bias–Variance $\text{Error} = \text{Bias}^2 + \text{Variance} + \text{Noise}$

Part V — Theory

- 41 Generalisation Bound $R \leq \hat{R} + \text{Complexity}$
- 42 PAC Learning $P(R(f) \leq \epsilon) \geq 1 - \delta$
- 43 Scaling Law $L(N) = A \cdot N^{-a}$
- 44 Power Law $y = ax^b$
- 45 Gaussian $N(\mu, \sigma^2)$
- 46 Expectation $E[X]$
- 47 Variance $\text{Var}(X) = E[(X - \mu)^2]$
- 48 Markov Property $P(S_{t+1} | S_t)$
- 49 Bellman Equation $V(s) = R + \gamma V(s')$
- 50 Universal Approximation $f(x) \approx \sum a_i \sigma(w_i^T x + b_i)$

How to read this

These fifty equations are not a syllabus to be memorised but a map of how the field thinks. They are arranged as five movements. Learning fixes what we are trying to do and the logic of doing it from data. Representation concerns how raw input becomes usable geometry. Computation is the machinery a model runs. Optimisation is how its parameters are actually set. Theory asks why any of it should work at all.

The same handful of ideas recur in disguise. A dot product becomes a neuron, an attention score, and a similarity. An entropy becomes a loss. An expectation becomes a risk. Reading across the entries, the point is to see the small number of moves the whole edifice is built from.

Part I — Learning

What it means to learn from data — the objects, the objectives, and the logic of updating belief.

01 Unknown Function

$$f(x)$$

Plain-English meaning. The thing you are trying to learn: a rule that turns inputs into outputs, whose exact form you do not know and must recover from data or approximate.

Mathematical intuition. A function is a deterministic map from inputs to outputs. “Unknown” means we treat it as a fixed but hidden object living in some space of functions, and learning is the search for a good stand-in.

Where it appears in AI. The conceptual starting point of all supervised learning. Image $\rightarrow$ label, sentence $\rightarrow$ next token, and state $\rightarrow$ action are all instances of an unknown f being approximated.

Historical origin. The $f(x)$ notation is Euler's (1734). Framing learning itself as the recovery of an unknown function is the founding move of statistical learning theory and function approximation.

Common misconceptions. That f must exist as a clean deterministic rule. Often there is no true function, only a conditional distribution, and the “unknown function” is a convenient idealisation.

Why it matters today. It fixes the target of the entire enterprise. Everything downstream — loss, risk, generalisation — is defined relative to this object we never fully observe.

02 Parameterised Function

$$f_{\theta}(x)$$

Plain-English meaning. Instead of searching over all possible functions, you pick a family of functions controlled by knobs (parameters θ) and learn the knob settings.

Mathematical intuition. θ indexes a manifold of candidate functions. Learning becomes optimisation over a finite-dimensional parameter space rather than the intractable space of all functions.

Where it appears in AI. Every trained model is an f_{θ} : the weights of a neural network, the coefficients of a regression. “Training” means choosing θ .

Historical origin. Parametric modelling is the backbone of Fisher's statistics in the early 20th century. The explicit f_{θ} framing is now standard and became central with neural networks.

Common misconceptions. That more parameters always means a richer or better model. Effective capacity depends on architecture and inductive bias, not raw parameter count.

Why it matters today. It is the bridge that makes learning computable. Modern models have billions of θ , and the entire training stack exists to set them.

03 Input–Output Mapping

$$f: X \rightarrow Y$$

Plain-English meaning. A formal statement that the function takes things from an input set X and returns things in an output set Y .

Mathematical intuition. It specifies the types — the domain and codomain — before any formula. Classification has discrete Y , regression continuous Y , generation structured Y .

Where it appears in AI. It defines the shape of the task: X = images, Y = labels; X = token sequences, Y = token distributions. The type of Y dictates loss and architecture.

Historical origin. The set-theoretic definition of a function as a map between sets crystallised with Dirichlet and the later Bourbaki formalisation of the 19th and 20th centuries.

Common misconceptions. That X and Y are simple. In practice they are high-dimensional, structured spaces — sequences, graphs, images — whose geometry is the hard part.

Why it matters today. Naming the domain and codomain is the first design decision in any system. Many failures are really mismatches between the assumed $X \rightarrow Y$ and reality.

04 Empirical Risk

$$\hat{R}(f) = (1/n) \sum L(y_i, f(x_i))$$

Plain-English meaning. The average mistake your model makes on the data you actually have.

Mathematical intuition. It is a Monte Carlo estimate of the true error using the sample — replacing an unknowable expectation with a computable average over n examples.

Where it appears in AI. This is literally the training loss you minimise. Fitting a model is empirical risk minimisation (ERM).

Historical origin. Formalised by Vapnik and Chervonenkis in statistical learning theory (1960s–70s) as the empirical counterpart to expected risk.

Common misconceptions. That low empirical risk means a good model. You can drive it to zero by memorising, which is exactly overfitting.

Why it matters today. It is the only error we can actually compute, so the whole theory of generalisation is about when minimising it is safe.

05 Expected Risk

$$R(f) = E[L(Y, f(X))]$$

Plain-English meaning. The average mistake your model would make over all possible data, not just your sample — the error you truly care about.

Mathematical intuition. An expectation of the loss under the true data distribution. It is the population quantity that empirical risk estimates.

Where it appears in AI. The real objective of learning. Test performance is an estimate of expected risk; everything else is a proxy for it.

Historical origin. Central to statistical decision theory (Wald) and statistical learning theory (Vapnik). The risk-as-expected-loss framing dates to mid-20th-century statistics.

Common misconceptions. That it can be measured directly. It depends on the unknown true distribution and can only ever be estimated on held-out data.

Why it matters today. It is the north star of learning; the gap between it and empirical risk is the entire subject of generalisation.

o6 Learning Objective

$$\theta^* = \arg \min L(\theta)$$

Plain-English meaning. Learning is choosing the parameters that make the loss as small as possible.

Mathematical intuition. It recasts learning as an optimisation problem: find the point in parameter space at the bottom of the loss surface.

Where it appears in AI. Every training run is an attempt to approximate this arg min via gradient descent. “Training” and “optimising $L(\theta)$ ” are synonyms.

Historical origin. The arg-min formulation is the language of mathematical optimisation (Fermat, Cauchy). Its dominance in ML grew with the optimisation-centric view of the 1980s–2010s.

Common misconceptions. That we find the true global minimum. In deep learning we find good, often flat, local minima; the exact optimum is neither reached nor even desirable.

Why it matters today. It unifies wildly different models under one recipe — define a loss, then descend — which is why the same optimiser trains vision, language, and control systems.

o7 Generalisation Gap

$$R(f) - \hat{R}(f)$$

Plain-English meaning. The difference between how well a model does on new data versus the data it trained on.

Mathematical intuition. The estimation error of empirical risk as a proxy for expected risk. Classically it shrinks with more data and grows with model capacity.

Where it appears in AI. The quantity behind train/test gaps, overfitting, and the need for validation sets. Controlling it is the practical goal of regularisation.

Historical origin. Made precise by VC theory (Vapnik–Chervonenkis) and later by Rademacher-complexity and PAC frameworks.

Common misconceptions. That a large gap always means failure. Modern over-parameterised models can generalise well despite classical bounds predicting disaster — the “double descent” puzzle.

Why it matters today. It is the core scientific question of ML: why fitting the seen data ever tells us anything about the unseen.

o8 Maximum Likelihood

$$\theta^* = \arg \max P(D|\theta)$$

Plain-English meaning. Choose the parameters under which the data you observed would have been most probable.

Mathematical intuition. Treat the data as fixed and the likelihood as a function of θ , then climb it to its peak. Taking logs turns products into sums and recovers cross-entropy.

Where it appears in AI. Training language models, classifiers, and most probabilistic models is maximum likelihood. Next-token prediction is MLE of a sequence model.

Historical origin. Introduced by R. A. Fisher (1912, formalised 1922), one of the pillars of modern statistics.

Common misconceptions. That MLE is always well-behaved. It can overfit badly with few samples and ignores prior information entirely.

Why it matters today. It is the statistical justification for the cross-entropy losses that train nearly every modern model.

09 Bayes Rule

$$P(A/B) = P(B/A)P(A)/P(B)$$

Plain-English meaning. A recipe for updating what you believe about a cause after seeing evidence.

Mathematical intuition. It inverts a conditional: it converts “probability of evidence given hypothesis” into “probability of hypothesis given evidence,” reweighted by prior plausibility.

Where it appears in AI. The foundation of Bayesian inference, probabilistic graphical models, Naive Bayes, Bayesian deep learning, and belief updating in agents.

Historical origin. Thomas Bayes' essay, published posthumously in 1763 by Richard Price; independently developed and generalised by Laplace.

Common misconceptions. That relying on priors makes it “unscientific.” The rule itself is a theorem of probability, entirely mechanical given its inputs.

Why it matters today. It is the canonical formalism for reasoning under uncertainty and updating beliefs — the logic beneath principled learning from evidence.

10 Posterior

$$P(\theta/D) \propto P(D/\theta)P(\theta)$$

Plain-English meaning. Your updated belief about the parameters is the data's evidence times your prior belief, up to a normalising constant.

Mathematical intuition. Posterior $\propto$ likelihood $\times$ prior. The “ $\propto$ ” hides the intractable evidence term $P(D)$, which is why Bayesian computation is hard.

Where it appears in AI. Bayesian neural networks, variational inference, Gaussian processes, and uncertainty quantification all revolve around approximating this posterior.

Historical origin. The Bayesian updating scheme of Bayes and Laplace, revived for ML through the 1990s–2000s (MacKay, Neal).

Common misconceptions. That the posterior is easy to compute. The normalising integral is usually intractable, forcing approximations such as MCMC or variational methods.

Why it matters today. It is the principled route to models that know what they don't know — increasingly vital as AI is deployed in high-stakes settings.

Part II — Representation

How raw inputs become useful internal form — geometry, similarity, and the information a representation should keep or discard.

11 Representation Function

$$\varphi : X \rightarrow Z$$

Plain-English meaning. A map that turns raw input into a more useful internal form — features the rest of the model can work with.

Mathematical intuition. φ re-coordinatises the input into a space Z where structure such as class or similarity becomes geometrically simple, ideally linearly separable.

Where it appears in AI. Embeddings, feature extractors, and the hidden layers of a network are all φ . Representation learning is the search for a good φ .

Historical origin. From hand-designed features (SIFT, TF-IDF) to learned features. The modern view that φ should be learned is the essence of deep learning.

Common misconceptions. That representations are neutral. Every φ encodes choices and biases and discards information deliberately.

Why it matters today. The quality of φ largely determines what a model can do. “Representation is everything” is the deep-learning thesis.

12 Latent Representation

$$z = \varphi(x)$$

Plain-English meaning. The compressed internal code an input gets mapped to inside a model.

Mathematical intuition. z is a point in a latent space. Nearby z 's should mean similar inputs, and directions in z -space often correspond to meaningful factors of variation.

Where it appears in AI. Embedding vectors, autoencoder bottlenecks, and the latent spaces of diffusion and other generative models.

Historical origin. Latent-variable modelling in statistics (factor analysis) evolved into learned latents via autoencoders and VAEs (Kingma & Welling, 2013).

Common misconceptions. That latent dimensions are individually interpretable. Meaning is usually distributed and entangled across many dimensions.

Why it matters today. Latent spaces are where modern generation, retrieval, and semantic search happen — computation over meaning rather than raw pixels or tokens.

13 Euclidean Distance

$$d(x,y) = \|x - y\|$$

Plain-English meaning. The straight-line distance between two points.

Mathematical intuition. The L2 norm of the difference vector — the length of the arrow connecting two points in space.

Where it appears in AI. k-nearest-neighbours, clustering such as k-means, and any similarity computation in embedding space.

Historical origin. Euclid's geometry, later formalised as a norm in the development of metric spaces and linear algebra.

Common misconceptions. That it is always the right notion of similarity. In high dimensions distances concentrate and Euclidean distance loses discriminative power.

Why it matters today. It is the default geometry of embedding spaces, though cosine similarity often supplants it when direction, not magnitude, carries meaning.

14 Cosine Similarity

$$(x \cdot y) / (\|x\| \|y\|)$$

Plain-English meaning. How aligned two vectors are in direction, ignoring their length.

Mathematical intuition. The cosine of the angle between vectors: +1 means the same direction, 0 orthogonal, -1 opposite. It is a magnitude-invariant similarity.

Where it appears in AI. The workhorse of semantic search, retrieval-augmented generation, and comparing text or word embeddings.

Historical origin. The vector space model of information retrieval, developed by Gerard Salton in the 1960s–70s.

Common misconceptions. That it measures distance. It measures angle, so two vectors of very different magnitude can still be maximally “similar.”

Why it matters today. Because meaning in embedding spaces is largely encoded by direction, cosine similarity is the default relevance score across modern retrieval and RAG systems.

15 Dot Product

$$x^T y$$

Plain-English meaning. Multiply two vectors component-wise and add up the results — a single number measuring their agreement.

Mathematical intuition. It combines magnitude and alignment; geometrically it equals $\|x\| \|y\| \cos \theta$. It is the atomic operation of linear algebra.

Where it appears in AI. Every neuron, attention score, and similarity is built from dot products. Matrix multiplication is dot products at scale.

Historical origin. Emerged from Grassmann's and Gibbs's vector algebra in the 19th century.

Common misconceptions. That it is a minor primitive. It is the single most executed operation in modern AI hardware, which is built to do it fast.

Why it matters today. GPUs and TPUs are essentially dot-product engines. The economics of AI rest on doing this operation trillions of times per second.

16 Principal Components (SVD)

$$X = USV^T$$

Plain-English meaning. Any dataset-matrix can be broken into its fundamental patterns, ranked by how much variation each explains.

Mathematical intuition. SVD factorises X into rotations (U , V^T) and a scaling (S). The top singular directions capture the dominant structure — the basis of PCA.

Where it appears in AI. Dimensionality reduction, PCA, latent semantic analysis, low-rank adaptation (LoRA), and compression.

Historical origin. SVD is due to Beltrami and Jordan (1870s); PCA to Pearson (1901) and Hotelling (1933).

Common misconceptions. That principal components are the “true” features. They are variance-maximising directions, which need not align with what is meaningful or causal.

Why it matters today. Low-rank structure underlies efficient fine-tuning (LoRA), model compression, and much of what makes large models tractable.

17 Information Bottleneck

$$\max I(Z;Y) - \beta I(Z;X)$$

Plain-English meaning. Squeeze the input through a bottleneck that keeps what predicts the target and throws away everything else.

Mathematical intuition. Trade off two mutual informations: keep Z informative about Y while forgetting the irrelevant detail of X , with β tuning the balance.

Where it appears in AI. A theoretical lens on what deep networks do — compress inputs into task-relevant representations — and a motivation for representation-learning objectives.

Historical origin. Introduced by Tishby, Pereira, and Bialek (1999); later proposed by Tishby as a theory of deep learning.

Common misconceptions. That it is an established explanation of deep learning. The “information bottleneck theory of deep learning” remains empirically contested.

Why it matters today. It offers a principled definition of a good representation — a minimal sufficient statistic — that continues to shape representation-learning research.

18 Entropy

$$H(X) = -\sum p(x) \log p(x)$$

Plain-English meaning. A measure of how uncertain or surprising a random quantity is on average.

Mathematical intuition. The expected number of bits needed to encode outcomes: maximal for uniform distributions, zero for certainties.

Where it appears in AI. The basis of cross-entropy loss, information gain in decision trees, and exploration bonuses in reinforcement learning.

Historical origin. Claude Shannon (1948), borrowing the mathematical form from Boltzmann–Gibbs statistical mechanics.

Common misconceptions. That entropy is “disorder.” In information theory it precisely quantifies uncertainty and compressibility, not physical mess.

Why it matters today. It is the currency of information. Nearly every training loss for classification and language is an entropy in disguise.

19 Mutual Information

$$I(X;Y) = H(X) - H(X|Y)$$

Plain-English meaning. How much knowing one variable reduces your uncertainty about another.

Mathematical intuition. The information shared between X and Y. It is zero if and only if they are independent, and it is symmetric in its arguments.

Where it appears in AI. Feature selection, representation learning (InfoMax and contrastive-learning bounds), and the information bottleneck.

Historical origin. Shannon (1948), as part of the founding of information theory.

Common misconceptions. That it captures only linear dependence. Unlike correlation it captures any dependence — but it is notoriously hard to estimate in high dimensions.

Why it matters today. Modern self-supervised methods, especially contrastive learning, are often justified as maximising a mutual-information lower bound.

20 KL Divergence

$$D_{KL}(P\|Q) = \sum P(x) \log(P(x)/Q(x))$$

Plain-English meaning. A measure of how far one probability distribution is from another — the cost of assuming Q when the truth is P.

Mathematical intuition. The expected extra bits from coding with the wrong distribution. It is always ≥ 0 , but asymmetric, so it is not a true distance.

Where it appears in AI. The objective behind variational inference (the ELBO), the regulariser in RLHF and PPO, model distillation, and VAEs.

Historical origin. Kullback and Leibler (1951), building on Shannon's information theory.

Common misconceptions. That it is a symmetric distance. $D_{KL}(P\|Q) \neq D_{KL}(Q\|P)$, and the direction changes behaviour — mode-covering versus mode-seeking.

Why it matters today. It is the mathematical leash that keeps fine-tuned models, such as RLHF policies, from drifting too far from a trusted reference.

Part III — Computation

The operations a model actually runs — the neuron, the layer, convolution, attention, and the tricks that make depth trainable.

21 Linear Neuron

$$y = w^T x + b$$

Plain-English meaning. Take a weighted sum of the inputs and add an offset — the simplest artificial neuron.

Mathematical intuition. An affine map, geometrically a hyperplane in input space. The weights set its orientation and the bias its threshold.

Where it appears in AI. The building block of every layer. A full network is many of these composed with nonlinearities.

Historical origin. The McCulloch–Pitts neuron (1943) and Rosenblatt's perceptron (1958).

Common misconceptions. That a stack of these alone is powerful. Without nonlinearities, any composition of linear neurons collapses to a single linear map.

Why it matters today. It is the irreducible unit of the entire field — everything scales up from this affine operation.

22 Activation

$$a = \sigma(w^T x + b)$$

Plain-English meaning. Pass the neuron's weighted sum through a nonlinear squashing function.

Mathematical intuition. The nonlinearity σ bends the space, letting stacked layers represent curved decision boundaries that linear maps cannot.

Where it appears in AI. Sigmoid, tanh, ReLU, GELU — the nonlinear heart of every deep network.

Historical origin. The logistic/sigmoid function dates to Verhulst (1838); its use as a neural activation grew with perceptrons and backpropagation.

Common misconceptions. That the specific activation barely matters. The shift from sigmoid to ReLU in the 2010s was pivotal for training deep networks by mitigating vanishing gradients.

Why it matters today. Nonlinearity is what separates deep learning from linear regression. Without σ there is no “deep.”

23 Layer

$$h = \sigma(Wx + b)$$

Plain-English meaning. A whole bank of neurons applied at once, transforming a vector of inputs into a vector of outputs.

Mathematical intuition. A matrix multiply (mixing) followed by an elementwise nonlinearity — the fundamental compute block, vectorised for hardware.

Where it appears in AI. The repeating unit of MLPs and, with variations, of every deep architecture.

Historical origin. The multilayer perceptron, made trainable by backpropagation (Rumelhart, Hinton, and Williams, 1986).

Common misconceptions. That width and depth are interchangeable. They trade off differently, and depth enables compositional feature hierarchies that width cannot cheaply replicate.

Why it matters today. Stacking these layers is the literal meaning of “deep” learning and the source of hierarchical representation.

24 Composition

$$f = g \circ h$$

Plain-English meaning. Apply one function, then feed its output into another.

Mathematical intuition. Depth is composition. Complex functions are built by chaining simple ones, and the chain rule makes this differentiable end to end.

Where it appears in AI. Every deep network is a long composition of layers. The “deep” in deep learning is compositional depth.

Historical origin. Function composition is elementary mathematics; its centrality to expressive power is a deep-learning-era insight.

Common misconceptions. That composing simple functions yields only simple behaviour. Composition is precisely where expressive, hierarchical power comes from.

Why it matters today. Compositionality is why deep networks reuse features and generalise — and why the chain rule (backpropagation) is the engine of training.

25 Matrix Multiplication

$$Y = WX$$

Plain-English meaning. Transform a whole batch of vectors at once by a linear operation.

Mathematical intuition. Each output entry is a dot product; the whole operation simultaneously rotates, scales, and mixes coordinates. It is linear maps made computational.

Where it appears in AI. The dominant operation in every forward and backward pass. Training cost is essentially counted in matmuls (FLOPs).

Historical origin. Matrix algebra formalised by Cayley (1858).

Common misconceptions. That it is just bookkeeping. The hardware, energy, and cost of modern AI are overwhelmingly determined by matmul throughput.

Why it matters today. The entire GPU and TPU industry exists to do this operation faster. It is the physical substrate of AI scaling.

26 Convolution

$$(f * g)(x)$$

Plain-English meaning. Slide a small pattern-detector across an input and record where it matches.

Mathematical intuition. A weighted local average that is translation-equivariant and weight-sharing — detect the same feature anywhere in the input.

Where it appears in AI. Convolutional neural networks for vision, audio, and signals; the inductive bias of locality and translation invariance.

Historical origin. Convolution is classical analysis (the Fourier era); CNNs trace to Fukushima's Neocognitron (1980) and LeCun's LeNet (1989).

Common misconceptions. That convolution is obsolete after Transformers. It remains dominant where locality and efficiency matter, and hybrid architectures persist.

Why it matters today. It encoded the first great inductive bias of deep learning and still powers efficient perception systems.

27 Attention Scores

$$QK^T$$

Plain-English meaning. Measure how much each element should pay attention to every other element.

Mathematical intuition. Dot products between query and key vectors produce a compatibility matrix — a data-dependent, all-pairs relevance map.

Where it appears in AI. The core of Transformers; how tokens gather context from other tokens.

Historical origin. Attention was introduced by Bahdanau et al. (2014) for translation; the QK^T self-attention form by Vaswani et al. (2017).

Common misconceptions. That attention weights are faithful explanations of a model's reasoning. Interpreting them as explanations is contested.

Why it matters today. This all-pairs interaction is what let models capture long-range dependencies and scale into today's large language models.

28 Scaled Dot-Product Attention

$$\text{softmax}((QK^T)/\sqrt{d})V$$

Plain-English meaning. Turn relevance scores into weights and use them to blend information from across the sequence.

Mathematical intuition. Normalise QK^T by $\sqrt{d}$ to control variance, softmax it into a distribution, then take a weighted average of the value vectors V .

Where it appears in AI. The exact mechanism at the heart of every Transformer block.

Historical origin. Vaswani et al., “Attention Is All You Need” (2017).

Common misconceptions. That the $\sqrt{d}$ scaling is cosmetic. Without it, large dot products saturate the softmax and kill gradients; it is essential for stable training.

Why it matters today. This single formula underpins the generative-AI era. GPT, Claude, and modern multimodal models are stacks of it.

29 Softmax

$$P_i = e^{z_i} / \sum e^{z_j}$$

Plain-English meaning. Turn a vector of scores into probabilities that sum to one.

Mathematical intuition. Exponentiate to make everything positive, then normalise. It is a smooth, differentiable approximation to “argmax.”

Where it appears in AI. The output layer of classifiers, the attention weighting, and the sampling distribution for language generation.

Historical origin. The Boltzmann/Gibbs distribution of statistical physics; named “softmax” by Bridle (1989).

Common misconceptions. That softmax outputs are calibrated probabilities. They are often overconfident and require calibration to be trusted.

Why it matters today. It is the standard interface between real-valued scores and probabilistic decisions — and its temperature controls the randomness of generation.

30 Residual Connection

$$y = x + F(x)$$

Plain-English meaning. Let a layer learn only the change to its input, then add that change back on.

Mathematical intuition. F learns a residual, and the identity shortcut gives gradients a clean path backward, preventing them from vanishing in very deep stacks.

Where it appears in AI. ResNets and every Transformer block; the reason networks can be hundreds of layers deep.

Historical origin. He, Zhang, Ren, and Sun, “Deep Residual Learning” (ResNet, 2015).

Common misconceptions. That it is a minor engineering trick. It was the breakthrough that made extreme depth trainable at all.

Why it matters today. Without residual connections, the deep Transformers behind modern AI simply would not train.

Part IV — Optimisation

How parameters are actually set — gradients, descent, the losses that shape behaviour, and the penalties that curb overfitting.

31 Loss

$$L(y, \hat{y})$$

Plain-English meaning. A number saying how wrong a single prediction is.

Mathematical intuition. A function comparing truth y to prediction $\hat{y}$. Its shape encodes what “wrong” means and what the model is pushed to prioritise.

Where it appears in AI. The scalar that training minimises. The choice of loss defines the task — regression, classification, ranking.

Historical origin. The loss function is central to Wald's statistical decision theory (1939–1950).

Common misconceptions. That the loss is a neutral technicality. It embeds values and trade-offs, and a mis-specified loss produces misaligned behaviour.

Why it matters today. “You get what you optimise for.” The loss is where human intent is (often imperfectly) translated into a training signal.

32 Gradient

$$\nabla L$$

Plain-English meaning. The direction in which the loss increases fastest, and how steeply.

Mathematical intuition. The vector of partial derivatives. It points uphill, so its negative points toward lower loss — the local compass for optimisation.

Where it appears in AI. Computed by backpropagation for every parameter. It is the signal that drives all learning.

Historical origin. Vector calculus (Hamilton's nabla, Gibbs). The gradient concept is 19th-century.

Common misconceptions. That the gradient points toward the minimum. It gives only the locally steepest direction, which can be a poor global route.

Why it matters today. Everything a model learns is transmitted through this vector. No gradient, no learning in the dominant paradigm.

33 Gradient Descent

$$\theta \leftarrow \theta - \eta \nabla L$$

Plain-English meaning. Repeatedly nudge the parameters a little in the downhill direction.

Mathematical intuition. Take small steps of size η against the gradient; the loss surface is descended iteratively toward a minimum.

Where it appears in AI. The core training loop of essentially all modern models, in stochastic, mini-batch form.

Historical origin. Cauchy (1847); the stochastic variant (Robbins–Monro, 1951) made it practical for large data.

Common misconceptions. That the learning rate η is a minor detail. It is often the single most important hyperparameter, and getting it wrong prevents learning entirely.

Why it matters today. This one update rule, scaled to billions of parameters and trillions of steps, is how every large model is trained.

34 Chain Rule

$$dz/dx = (dz/dy)(dy/dx)$$

Plain-English meaning. To find how a change ripples through a chain of functions, multiply the local rates of change.

Mathematical intuition. Derivatives compose multiplicatively. This is exactly how gradients flow backward through composed layers.

Where it appears in AI. The mathematical basis of backpropagation — the algorithm that makes deep learning possible.

Historical origin. Leibniz and the founding of calculus; its algorithmic use as backpropagation was popularised by Rumelhart, Hinton, and Williams (1986).

Common misconceptions. That backpropagation is a special AI invention. It is the chain rule applied systematically — reverse-mode automatic differentiation.

Why it matters today. Every gradient in every deep model is computed by the chain rule. It is the literal engine of training.

35 Mean Squared Error

$$(1/n) \sum (y - \hat{y})^2$$

Plain-English meaning. Average the squared gaps between predictions and truth.

Mathematical intuition. Squaring penalises large errors heavily and is smooth and differentiable. Minimising it targets the conditional mean.

Where it appears in AI. The default regression loss and a building block in many objectives, including diffusion-model training.

Historical origin. Least squares, due to Legendre (1805) and Gauss (1809).

Common misconceptions. That MSE is always appropriate. It is sensitive to outliers and implicitly assumes Gaussian noise.

Why it matters today. It is the canonical loss for continuous targets and, via its Gaussian link, connects optimisation back to maximum likelihood.

36 Cross Entropy

$$-\sum y \log \hat{y}$$

Plain-English meaning. Penalise a classifier for putting low probability on the correct answer.

Mathematical intuition. It is the KL divergence between the true and predicted label distributions (up to a constant). Minimising it is maximum likelihood for classification.

Where it appears in AI. The standard loss for classification and for next-token prediction in language models.

Historical origin. Derived from Shannon's entropy (1948); standard in ML classification from the 1990s onward.

Common misconceptions. That it and MSE are interchangeable for classification. Cross-entropy gives far better gradients for probabilistic outputs.

Why it matters today. It is the loss that trains virtually every large language model — pretraining is cross-entropy on next tokens.

37 L2 Regularisation

$$L + \lambda \|w\|^2$$

Plain-English meaning. Add a penalty for large weights to keep the model simpler.

Mathematical intuition. It shrinks weights smoothly toward zero (weight decay), preferring small, spread-out weights and improving conditioning.

Where it appears in AI. Weight decay in essentially every trained network; ridge regression in classical ML.

Historical origin. Tikhonov regularisation (1943/1963); ridge regression by Hoerl and Kennard (1970).

Common misconceptions. That weight decay and the L2 penalty are identical under all optimisers. With adaptive methods they differ — hence AdamW.

Why it matters today. It is the most common defence against overfitting and a standard ingredient of stable large-scale training.

38 L1 Regularisation

$$L + \lambda \|w\|_1$$

Plain-English meaning. Penalise the sum of weight magnitudes to push many weights exactly to zero.

Mathematical intuition. The L1 penalty's corners drive sparsity — it selects features by zeroing out the rest.

Where it appears in AI. Sparse models, feature selection, compressed sensing, and interpretable linear models.

Historical origin. The LASSO, introduced by Tibshirani (1996).

Common misconceptions. That L1 and L2 both just “shrink.” Only L1 produces true sparsity (exact zeros), which is qualitatively different.

Why it matters today. Sparsity aids interpretability and efficiency — themes resurgent in model compression and mechanistic interpretability.

39 Adam Update

$$\theta \leftarrow \theta - \eta \hat{m} / (\sqrt{\hat{v}} + \epsilon)$$

Plain-English meaning. An optimiser that adapts each parameter's step size using recent gradient history.

Mathematical intuition. Combine momentum (a running mean $\hat{m}$ of gradients) with per-parameter scaling by gradient variance ($\hat{v}$): bigger steps where gradients are small and steady.

Where it appears in AI. The default optimiser for training most deep networks, including large language models (usually the AdamW variant).

Historical origin. Kingma and Ba (2014).

Common misconceptions. That Adam always beats plain SGD. SGD with momentum often generalises better in vision, and the choice is task-dependent.

Why it matters today. Adam's robustness to hyperparameters is a large part of why training huge, heterogeneous models became practical.

40 Bias–Variance

$$\text{Error} = \text{Bias}^2 + \text{Variance} + \text{Noise}$$

Plain-English meaning. A model's error splits into being systematically off, being too sensitive to the data, and irreducible randomness.

Mathematical intuition. A decomposition of expected error. Classically, bias falls and variance rises as model complexity grows — a U-shaped tradeoff.

Where it appears in AI. The conceptual frame for underfitting versus overfitting and for choosing model capacity.

Historical origin. A classical statistics result, sharpened for neural networks by Geman, Bienenstock, and Doursat (1992).

Common misconceptions. That the tradeoff is universal. “Double descent” shows very large models can lower both, complicating the classical U-curve.

Why it matters today. It is the intuition pump for generalisation, and its apparent breakdown in the deep-learning regime is an active research frontier.

Part V — Theory

Why any of it should work — bounds, learnability, the statistics of scale, and the mathematics of sequential decision-making.

41 Generalisation Bound

$$R \leq \hat{R} + \text{Complexity}$$

Plain-English meaning. True error is at most training error plus a penalty for how complex the model is.

Mathematical intuition. A high-probability upper bound. The complexity term — VC dimension or Rademacher complexity — captures how much the model could overfit.

Where it appears in AI. The theoretical backbone of statistical learning theory and the justification for regularisation.

Historical origin. Vapnik–Chervonenkis theory (1970s) and later Rademacher-complexity bounds.

Common misconceptions. That these bounds are tight for deep nets. Classical bounds are famously vacuous for over-parameterised models — a major open problem.

Why it matters today. It frames the central theoretical question: why do models that could memorise anything nonetheless generalise?

42 PAC Learning

$$P(R(f) \leq \epsilon) \geq 1 - \delta$$

Plain-English meaning. With high probability ($1 - \delta$), the learned model's error is small ($\leq \epsilon$), given enough data.

Mathematical intuition. “Probably Approximately Correct” formalises learnability as achieving low error with high confidence using a polynomial number of samples.

Where it appears in AI. The foundational framework of computational learning theory; it defines what it means for a hypothesis class to be learnable.

Historical origin. Leslie Valiant (1984), work that later won him the Turing Award.

Common misconceptions. That PAC guarantees describe deep-learning practice. The framework is worst-case and distribution-free, often loose for real models.

Why it matters today. It gave learning a rigorous computational definition, setting the terms for decades of theory.

43 Scaling Law

$$L(N) = A \cdot N^{-\alpha}$$

Plain-English meaning. Loss falls predictably as a power of model or data size — bigger is smoothly, quantifiably better.

Mathematical intuition. A straight line on a log-log plot. The exponent α tells you the diminishing but reliable returns to scale.

Where it appears in AI. Predicting and planning the training of frontier models, and deciding how to trade off parameters, data, and compute.

Historical origin. Hestness et al. (2017); Kaplan et al. (2020); refined by Hoffmann et al. (Chinchilla, 2022).

Common misconceptions. That scaling laws promise indefinite improvement. They describe smooth loss trends, not capability guarantees, and can bend or break.

Why it matters today. These empirical laws turned model-building into a forecastable engineering discipline and drove the compute race of the 2020s.

44 Power Law

$$y = ax^b$$

Plain-English meaning. One quantity grows or shrinks as a power of another — the shape behind “the rich get richer.”

Mathematical intuition. A scale-invariant relationship: linear on log-log axes, with heavy tails and no characteristic scale.

Where it appears in AI. Underlies scaling laws, token-frequency distributions (Zipf), and the statistics of natural data.

Historical origin. Pareto (income, 1890s) and Zipf (word frequencies, 1930s–40s); ubiquitous across complex systems.

Common misconceptions. That anything roughly straight on log-log is a power law. Genuine power laws require careful statistical validation.

Why it matters today. The power-law structure of language and data is part of why scaling works and why rare events dominate model behaviour.

45 Gaussian

$$N(\mu, \sigma^2)$$

Plain-English meaning. The classic bell curve — data clustered around an average with a characteristic spread.

Mathematical intuition. Fully described by mean and variance. It is the maximum-entropy distribution for a fixed variance and the limit of many summed effects.

Where it appears in AI. Weight initialisation, noise models, VAEs, diffusion models, and Gaussian processes.

Historical origin. De Moivre (1733), Gauss (1809), and Laplace; the Central Limit Theorem explains its ubiquity.

Common misconceptions. That real data is Gaussian. Much real-world data is heavy-tailed, and assuming normality can be badly wrong.

Why it matters today. It is the default noise and prior across ML, and the literal engine of diffusion generative models, which add and remove Gaussian noise.

46 Expectation

$$E[X]$$

Plain-English meaning. The long-run average value of a random quantity.

Mathematical intuition. A probability-weighted average over outcomes — the linear operator at the centre of probability and statistics.

Where it appears in AI. Expected risk, expected reward in RL, and the target of virtually every objective, which are expectations we estimate by sampling.

Historical origin. Huygens (1657), from the Pascal–Fermat correspondence on games of chance.

Common misconceptions. That the expectation is a typical or likely value. For skewed distributions it can be a value that rarely or never occurs.

Why it matters today. Nearly every loss and objective in ML is an expectation, and training is the art of estimating and minimising it from samples.

47 Variance

$$\text{Var}(X) = E[(X - \mu)^2]$$

Plain-English meaning. How spread out a random quantity is around its average.

Mathematical intuition. The expected squared deviation from the mean. Its square root, the standard deviation, is in the original units.

Where it appears in AI. Uncertainty estimates, the variance term in bias–variance, gradient-noise analysis, and normalisation layers.

Historical origin. The term “variance” was coined by R. A. Fisher (1918).

Common misconceptions. That low variance is always good. Some variance — exploration, stochasticity — is essential for learning and generalisation.

Why it matters today. Controlling variance of gradients, activations, and predictions is a recurring theme, from batch-norm to ensembling to uncertainty quantification.

48 Markov Property

$$P(S_{t+1} | S_t)$$

Plain-English meaning. The future depends only on the present state, not on the full history of how you got there.

Mathematical intuition. Memorylessness: the current state is a sufficient statistic for predicting the next one.

Where it appears in AI. The foundation of Markov Decision Processes (the model of reinforcement learning) and Markov chains, including MCMC.

Historical origin. Andrey Markov (1906), studying dependent sequences.

Common misconceptions. That real problems are truly Markov. Often they are not, and “state” must be engineered or learned so the property approximately holds.

Why it matters today. It is the modelling assumption that makes sequential decision-making and reinforcement learning tractable.

49 Bellman Equation

$$V(s) = R + \gamma V(s')$$

Plain-English meaning. The value of a state is the immediate reward plus the discounted value of where you land next.

Mathematical intuition. A recursive self-consistency condition on value. Solving it yields optimal long-term behaviour, and γ trades off now versus later.

Where it appears in AI. The heart of reinforcement learning — value iteration, Q-learning, and deep RL all descend from it.

Historical origin. Richard Bellman (1957) and the theory of dynamic programming.

Common misconceptions. That it requires knowing the environment's dynamics. Model-free methods such as Q-learning solve it from experience alone.

Why it matters today. It is the mathematical core of agents that plan and learn from reward, from game-playing systems to RLHF.

50 Universal Approximation

$$f(x) \approx \sum a_i \sigma(w_i^T x + b_i)$$

Plain-English meaning. A neural network with one hidden layer can approximate essentially any continuous function, given enough neurons.

Mathematical intuition. Weighted sums of nonlinear “bumps” can build up any well-behaved shape to arbitrary accuracy on a bounded domain.

Where it appears in AI. The existence proof that neural networks are, in principle, expressive enough for any task.

Historical origin. Cybenko (1989) for sigmoids; Hornik, Stinchcombe, and White (1989).

Common misconceptions. That it explains deep learning's success. It guarantees existence, not learnability or efficiency, is silent on why depth helps, and does not say gradient descent will find the solution.

Why it matters today. It underwrites the field's confidence that networks can represent what we need — while its gaps, on why depth and why trainable, frame what theory still owes us.