Enterprise Context Graphs

Data + Decision as the Foundation for Institutional AI Agents

Ajit Jaokar | Erdos Research

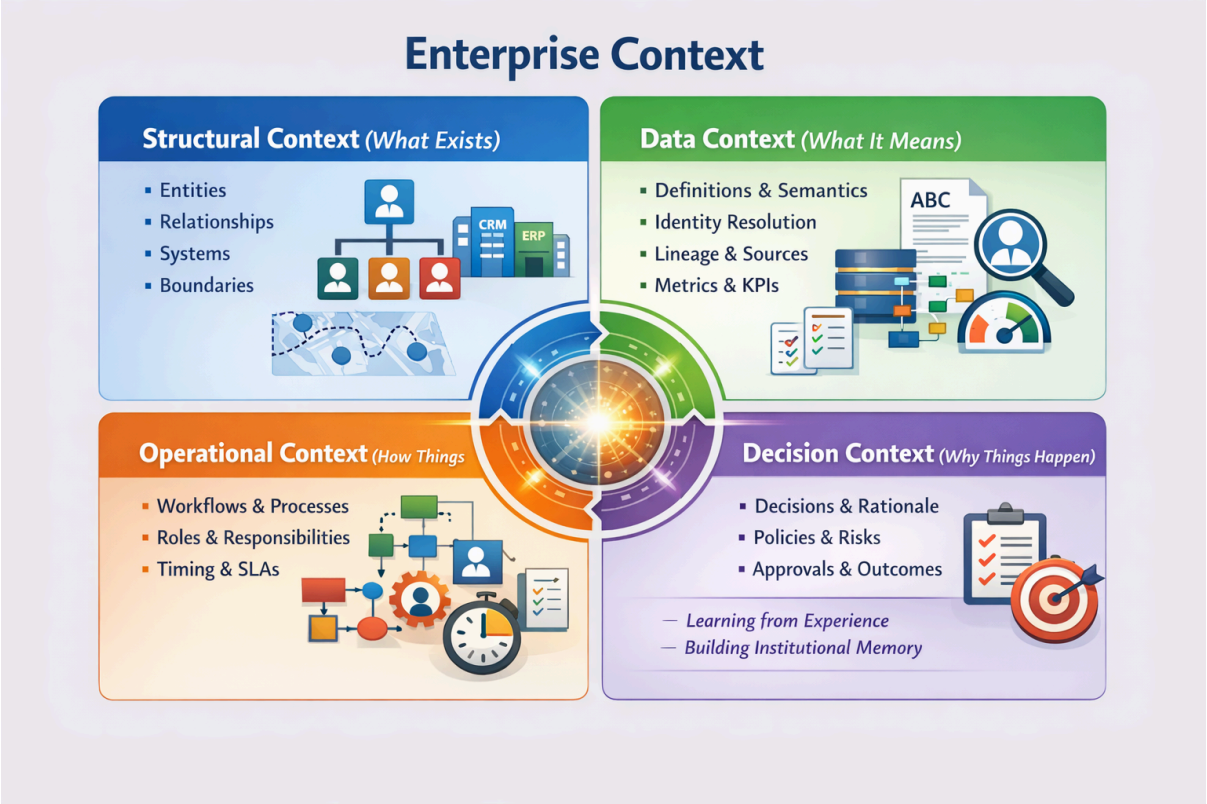

Introduction

The enterprise AI landscape is awash with agents. Every major platform vendor now offers an agentic framework, and early adopters are piloting autonomous workflows across functions from procurement to customer support. Yet a striking pattern is emerging: agents that perform well in controlled demos routinely fail when they encounter the full complexity of real organisational environments. They hallucinate policy, misinterpret metrics, act out of sequence, and make decisions that no informed employee would make.

The root cause is not a deficit of reasoning capability. It is a deficit of **context**.

Most agent architectures treat the enterprise as a flat retrieval surface: a collection of documents, database rows and API endpoints to be searched on demand. This works for answering questions, but it is fundamentally inadequate for taking action. Action requires understanding what the organisation *is*, what it *knows*, how it *operates*, and why it *acts*. In short, action requires enterprise context.

This article argues that enterprise context—modelled as a multi-layered context graph encompassing data context and decision context—is the necessary foundation for AI agents that can act correctly, consistently and institutionally within an organisation. We define the concept, decompose it into four interlinked layers, and show how each layer maps to the stages of the agentic reasoning loop.

Defining Enterprise Context

Enterprise context is a multi-layered, machine-interpretable representation of an organisation's identity and intelligence. Specifically, it captures:

- **What the enterprise is** — its entities, structures and systems.
- **What the enterprise knows** — its data, documents and semantics.
- **How the enterprise operates** — its workflows, constraints and dependencies.
- **Why the enterprise acts** — its decisions, rationale, policies and precedents.

Because our primary interest is in AI agents, we can sharpen this into a working definition:

> *Enterprise context is the minimum complete representation of an organisation required for an agent to act correctly, consistently and institutionally within it.*

The phrase "minimum complete" is deliberate. It forces us to include not just data, but **meaning** and **judgement** as foundations for action. A database of customer records is data. Knowing which definition of "active customer" applies in a given business unit is meaning. Knowing why that definition was chosen over alternatives—and what happened the last time someone applied the wrong one—is judgement. All three must be present for an agent to operate at an institutional level.

We can refine this further into a formal statement:

> *Enterprise context is a continuously evolving, context graph-structured representation of an organisation's entities, meanings, processes and decisions—enabling agents to act with accuracy, consistency and institutional judgement.*

The Four Layers of Enterprise Context

Rather than treating context as a monolithic concept, it is productive to decompose it into four tightly coupled layers. Each serves a distinct function; together, they provide the complete substrate an agent needs to reason and act.

Layer 1: Structural Context — What Exists

Structural context is the ontology of the enterprise. It defines the entities the organisation recognises (customers, suppliers, contracts, employees, assets), the relationships between them (owns, supplies, approves, depends_on), the systems that manage them (ERP, CRM, HR, finance, ticketing), and the boundaries that partition them (regions, business units, product lines).

This is the layer where relational systems have traditionally dominated. It stores the operational truth of the enterprise: the canonical record of what exists and how it is connected. For an agent, structural context answers the question: *What is the landscape I am operating in?*

Layer 2: Data Context — What It Means

Data context is the semantic layer. Where structural context tells the agent what exists, data context tells it what those things *mean*. It resolves definitions ("What is revenue?" "What counts as an active customer?"), reconciles identity across systems (the same supplier appearing under different names in different platforms), traces lineage (where a data point originated and how it was transformed), establishes source-of-truth hierarchies (which system's value wins in a conflict), and codifies metric logic (how KPIs are computed).

This is the layer that makes an agent **accurate and consistent**. Without it, the agent is guessing—and in an enterprise environment, guessing carries real cost.

Layer 3: Operational Context — How Things Happen

Operational context captures the dynamics of the enterprise: its workflows (procurement, claims processing, onboarding, incident response), the states and transitions that govern them, the roles and responsibilities assigned at each step, the timing and sequencing constraints, and the service-level agreements and compliance rules that bound permissible action.

This layer is conspicuously absent from most "data-only" approaches to enterprise AI. Yet it answers two questions that are indispensable for action: *Where in the process are we?* and *What is allowed to happen next?*

Operational context is what makes an agent **operationally correct**—not merely analytically correct. An agent that identifies the right answer but acts out of sequence, skips an approval gate, or violates an SLA has not acted correctly in any meaningful institutional sense.

Layer 4: Decision Context — Why Things Happen

Decision context is the most important and the most differentiated layer. It captures the decisions the organisation has made, the alternatives that were considered, the evidence that was weighed, the policies that were applied, the risk thresholds that governed the outcome, the approvals that were obtained, and the results that followed.

This is the layer that makes an agent **institutional**. It enables reasoning from precedent, consistent judgement across cases, and alignment with organisational norms. An agent grounded in decision context does not merely retrieve a policy document and attempt to interpret it from first principles. It can see how that policy has been applied in analogous situations, what edge cases have arisen, and how the organisation has responded.

Critically, **this layer compounds over time**. Every decision recorded enriches the graph. Every outcome logged sharpens future reasoning. Over months and years, the decision context layer becomes a genuine institutional asset—a **strategic moat** that is difficult for competitors to replicate because it encodes the accumulated judgement of the organisation.

Grounding the Agent Loop in Enterprise Context

An AI agent, at its most abstract, is defined by a loop: **Reason** → **Act** → **Observe** → **Update**. The critical insight is that this loop does not operate in isolation. In an enterprise setting, it must be embedded inside enterprise context. The generic loop becomes a context-grounded loop:

> *Reason (with context) → Act (within context) → Observe (from context) → Update (the context)*

Each stage of the loop draws on specific layers of the enterprise context graph. The following table summarises the mapping.

Step	Context Layers Used	What the Agent Does
Reason	Data + Decision Context	Resolve entities, interpret metrics, retrieve precedents, identify policy constraints
Act	Operational + Structural Context	Follow workflow sequences, respect approval gates, interact with correct systems
Observe	Structural + Data Context	Read updated records, track workflow progression, register new relationships
Update	Decision Context (primary)	Store decision traces: rationale, evidence, action taken, outcome observed

Step 1: Reason — "Understand and Decide What to Do"

The agent receives a task, query or trigger alongside the current enterprise state. Before it begins reasoning, it grounds itself in context. Through the data context layer, it resolves which entities are involved, interprets what metrics and thresholds mean, and identifies the relevant relationships and history. Through the decision context layer, it retrieves similar past cases, understands the precedents and policies that apply, and learns the typical approval patterns.

This transforms reasoning from generic LLM inference into **institutional reasoning**—reasoning that reflects the accumulated knowledge and judgement of the organisation.

Step 2: Act — "Do Something in the Enterprise"

The agent may make a recommendation, trigger a workflow, call an API, update a system, or request human approval. The operational context layer tells it which step of the workflow it is in, what actions are permitted at this stage, who needs to approve, and what compliance or SLA constraints apply. The structural context layer tells it which systems to interact with and which entities are affected.

This ensures the agent does not act out of sequence, does not violate established processes, and behaves as a knowledgeable employee would—not as an external tool with API access.

Step 3: Observe — "See What Happened"

After acting, the agent collects new data, system responses, human feedback and outcome signals. Through the structural and data context layers, it reads updated records (an invoice marked as paid, a ticket closed, a risk score recalculated) and registers new relationships (a newly linked supplier). Through the operational context layer, it tracks workflow progression: was the action approved, rejected or escalated?

This grounds observation in **real enterprise state**, not merely in text feedback or output parsing.

Step 4: Update — "Learn and Record"

This is the most under-designed step in most agent systems—and, we argue, the most consequential. The agent records what it did, why it did it, what happened, and whether the outcome was correct. This is written primarily to the decision context layer as a **decision trace**: the context in which the decision was made, the rationale applied, the evidence considered, the action taken, and the outcome observed. Optionally, the data context layer may also be updated with derived features or new semantic links.

This is the step where **the enterprise learns**. Decision traces are not ephemeral logs. They are structured, queryable records that become available to future reasoning cycles. Over

time, they constitute a compounding asset: each decision makes the next one better informed.

From Agent Loop to Context-Aware Agent Loop

With the four-layer model in place, we can rewrite the agent loop more precisely:

> *Perceive → Ground → Reason → Act → Observe → Record → Improve*

In this expanded formulation, **Ground** means querying the data context to resolve entities and semantics; **Reason** means combining data and decision context to formulate a plan; **Act** means executing within the constraints of operational context; **Observe** means reading updated state from the structural and data layers; **Record** means writing decision traces to the decision context layer; and **Improve** means feeding recorded outcomes back into the reasoning substrate for future cycles.

The following table contrasts a typical agent with a context-aware enterprise agent across the key dimensions.

Dimension	Typical Agent	Context-Aware Agent
Reasoning	Prompt + retrieved documents	Structured enterprise context
Action	Via generic tools	Within workflow constraints
Observation	Output parsing	Real system state
Learning	Internal vector memory	Organisational decision memory
Institutional Alignment	Weak	Strong

The Core Thesis: Context + Loop = Institutional Intelligence

The real shift this framework articulates is simple but profound: **the agent loop is not the intelligence**. The context and the loop together create intelligence.

- **Context without a loop** is static knowledge—a rich but inert representation of the organisation.
- **A loop without context** is shallow intelligence—capable reasoning that lacks the grounding to be institutionally valid.

- **Context combined with a loop** is institutional intelligence—an agent that reasons, acts, learns and improves within the living model of the organisation.

Each of the four context layers contributes a distinct dimension of this intelligence:

- **Data context** makes reasoning correct.
- **Operational context** makes actions valid.
- **Structural context** makes interactions grounded.
- **Decision context** makes behaviour consistent and improving over time.

Implications for Enterprise AI Strategy

If this analysis is correct, it has several practical consequences for organisations building or adopting AI agents.

Invest in context before investing in agents

The most common failure mode in enterprise AI is deploying increasingly capable agents on top of impoverished context. Organisations that invest in building a rich, structured enterprise context graph—encompassing all four layers—will find that even relatively simple agents produce dramatically better outcomes than sophisticated agents operating without context.

Treat decision context as a strategic asset

Most organisations have invested heavily in data infrastructure (the structural and data context layers) but have invested almost nothing in systematically capturing decisions, rationale and outcomes. This is a missed opportunity. Decision context is the layer that compounds, the layer that differentiates, and the layer that creates a genuine competitive moat. Organisations should begin capturing decision traces now, even before their agent capabilities are mature, because the value of this layer grows with every recorded decision.

Design agents as context-embedded systems

Agent architectures should be designed from the outset to operate within enterprise context, not bolted onto it after the fact. This means building explicit interfaces between the agent loop and each context layer: mechanisms for grounding, for constraint enforcement, for state observation, and above all for decision recording. The “update” step of the loop deserves as much engineering attention as the “reason” step.

Recognise the compounding dynamics

The enterprise context graph is not a static artefact. It is a living, evolving model that improves with every agent interaction. Organisations that understand this compounding dynamic will prioritise early adoption—not because their first agents will be perfect, but because every imperfect decision, once recorded and reviewed, makes the next one better.

The gap between organisations that begin building decision context early and those that delay will widen over time.

Conclusion

An enterprise AI agent is not just a reasoning loop. It is a context-embedded loop operating over a living model of the organisation. The context graph—spanning structure, data, operations and decisions—is what transforms generic AI capability into institutional intelligence: the ability to act correctly, consistently and with the accumulated judgement of the organisation.

The organisations that will lead in the age of agentic AI will not be those with the most powerful models or the most sophisticated tool chains. They will be those that build the deepest, richest enterprise context—and, above all, those that begin capturing and compounding decision context before their competitors recognise its value.